

# EX 03 - Data Handling & MLE

*Afek Adler*

*2020-03-22*

Last exercise we did:

- Expectency and Variance of the sample mean and sample sum
- Central limit theorem
- Bias variance decomposition of a point estimator
- Derived an unbiased estimate for  $\sigma^2(S^2)$
- Covered the student's t-distribution and chi square distribution

Today we will:

- Cover methods for point estimation
- Get to know `dplyr` package
- Try to develop a feeling for bayesian estimation.

## Loss function

A quick recap of the MSE of an estimator:

$$\text{MSE}(\hat{\theta}) = E((\hat{\theta} - \theta)^2)$$

The squared loss did not come from heaven but from convenience. for example, another good criterion can be:

$$\text{MAE}(\hat{\theta}) = E(|\hat{\theta} - \theta|)$$

Or many other types of error function. Also, at the lecture you have seen an *example* of Bayesian estimation where  $\hat{\theta}_{\text{MMSE}} = \int \theta p(\theta|\mathbf{x})d\theta = E(\theta|\mathbf{x})$ , the derivation of this formula was taken under assumption of a square loss but there are also many other bayesian estimators like the maximum a posteriori estimation -  $\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta|x)$ . In the end of the exercise we will go deeper into this subject.

## Point estimaion

### Some nice to have charactraists

- Unbiased. If  $E(\hat{\theta}) = \theta$
- Consistent. If the varaince of the estimator  $\sim 0$  when  $N$  tends to  $\infty$

Remember that the sample mean is unbiased estimator of the population mean

### Point estimates with the method of moments (MOM)

The first moment

$$E(X) = \frac{\sum_{i=1}^n X_i}{n} \Rightarrow E(X) = \bar{X}$$

The Second moment

$$E(X^2) = \frac{\sum_{i=1}^n X_i^2}{n} \Rightarrow$$

$$V(X) = E(X^2) - E^2(X) \Rightarrow V(X) = \frac{\sum_{i=1}^n X_i^2}{n} - (\bar{X})^2$$

### Q1: MOM

Let

$$X = \mathcal{U}(\theta, \theta + 6)$$

Estimate  $\theta$  with the method of moments

Therefore

$$E(X) = (\theta + \theta + 6)/2 = \theta + 3$$

And

$$E(X) = \bar{X} = \theta + 3$$

By the equation of the first moment. Therefore

$$\hat{\theta} = \bar{X} - 3$$

### Point estimates with the maximum likelihood estimation (MLE)

In statistics, maximum likelihood estimation (MLE) is a method of estimating the parameters of a probability distribution by maximizing a likelihood function, so that under the assumed statistical model the observed data is most probable. The point in the parameter space that maximizes the likelihood function is called the maximum likelihood estimate. The logic of maximum likelihood is both intuitive and flexible, and as such the method has become a dominant means of statistical inference.

If the likelihood function is differentiable, the derivative test for determining maxima can be applied. In some cases, the first-order conditions of the likelihood function can be solved explicitly; for instance, the ordinary least squares estimator maximizes the likelihood of the linear regression model. Under most circumstances, however, numerical methods will be necessary to find the maximum of the likelihood function. ("Wikipedia")

**Q2: MLE** With the binomial distribution - suppose we had a trial with 49 success out of 80.

$$L(p) = f_D(H = 49|p) = \binom{80}{49} p^{49} (1-p)^{31} \quad (1)$$

$$0 = \frac{\partial}{\partial p} \left( \binom{80}{49} p^{49} (1-p)^{31} \right), \{discard binomial coefficient\} \quad (2)$$

$$0 = 49p^{48}(1-p)^{31} - 31p^{49}(1-p)^{30}, \{(uv)' = u'v + v'u\} \quad (3)$$

$$= p^{48}(1-p)^{30}[49(1-p) - 31p] \quad (4)$$

$$= p^{48}(1-p)^{30}[49 - 80p] \quad (5)$$

Can be solved also by applying log on the likelihood.

It's clear that the maximum is at  $p = 49/80$ . But let's see how we do it in R using the built-in optimize function:

```
likelihood <- function(p) {  
  p^49*((1-p)^31)  
}  
tolerance <- 10^(-4)  
pmax <- optimize(likelihood, c(0, 1), tol = tolerance, maximum = T)[[1]]  
delta <- abs(pmax - (49/80))  
delta
```

```
## [1] 6.814623e-07
```

## HW1 q3

### Best Practices for data Data handling with R

R main datatypes:

- vectors
- matrices
- data.frame - matrices with metadata, added functionality and allow multiple data types
- tibbles - modern take on dataframes

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `mutate()` adds new variables that are functions of existing variables.
- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `summarize()` reduces multiple values down to a single summary.
- `arrange()` sorts the rows.

```
library(tidyverse)
library(nycflights13)
```

This dataset has 19 columns so the head function is not that useful when knitting to html. It is always useful to know how many missing values we have in our dataset, sometimes missing values are not just given to us as NA.

```
head(flights,2)
```

```
## # A tibble: 2 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517             515         2     830
## 2  2013     1     1     533             529         4     850
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
colSums(is.na(flights))/nrow(flights)
```

```
##           year           month           day           dep_time sched_dep_time
## 0.000000000 0.000000000 0.000000000 0.024511842 0.000000000
## dep_delay      arr_time sched_arr_time      arr_delay      carrier
## 0.024511842 0.025871796 0.000000000 0.028000808 0.000000000
## flight      tailnum      origin      dest      air_time
## 0.000000000 0.007458964 0.000000000 0.000000000 0.028000808
## distance      hour      minute      time_hour
## 0.000000000 0.000000000 0.000000000 0.000000000
```

```
sapply(flights,class)
```

```
## $year
## [1] "integer"
##
## $month
## [1] "integer"
##
```

```

## $day
## [1] "integer"
##
## $dep_time
## [1] "integer"
##
## $sched_dep_time
## [1] "integer"
##
## $dep_delay
## [1] "numeric"
##
## $arr_time
## [1] "integer"
##
## $sched_arr_time
## [1] "integer"
##
## $arr_delay
## [1] "numeric"
##
## $carrier
## [1] "character"
##
## $flight
## [1] "integer"
##
## $tailnum
## [1] "character"
##
## $origin
## [1] "character"
##
## $dest
## [1] "character"
##
## $air_time
## [1] "numeric"
##
## $distance
## [1] "numeric"
##
## $hour
## [1] "numeric"
##
## $minute
## [1] "numeric"
##
## $time_hour
## [1] "POSIXct" "POSIXt"

```

At home - find a better way to print the classes and the % of missing values in R

**select()** picks variables based on their names.

```
flight_distance_airtime <- flights %>% select( distance, air_time)
flight_distance_airtime %>% head(2)
```

```
## # A tibble: 2 x 2
##   distance air_time
##   <dbl>    <dbl>
## 1    1400      227
## 2    1416      227
```

**mutate()** adds new variables that are functions of existing variables.

```
flight_distance_airtime %>% mutate(mean_speed = distance/air_time) %>% head(2)
```

```
## # A tibble: 2 x 3
##   distance air_time mean_speed
##   <dbl>    <dbl>    <dbl>
## 1    1400      227      6.17
## 2    1416      227      6.24
```

If you only want to keep the new variables, use **transmute()**:

```
flight_distance_airtime %>% transmute(mean_speed = distance/air_time) %>% head(2)
```

```
## # A tibble: 2 x 1
##   mean_speed
##   <dbl>
## 1      6.17
## 2      6.24
```

**filter()** picks cases based on their values.

```
flights %>% filter(is.na(dep_delay)) %>% head(2)
```

```
## # A tibble: 2 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1      NA             1630         NA       NA
## 2  2013     1     1      NA             1935         NA       NA
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

**arrange()** picks cases based on their values.


```
flights %>% arrange(desc(month)) %>% head(2)
```

```
## # A tibble: 2 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013    12     1     13             2359         14     446
## 2  2013    12     1     17             2359         18     443
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```


```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

`summarize()` reduces multiple values down to a single summary.

```
by_month <- group_by(flights, month)
by_month %>% summarise(count = n()) %>%
  ggplot(mapping = aes(x = month, y = count)) + geom_bar(stat="identity") + coord_cartesian(ylim = c(
```



```
ggplot(data = flights) +
  geom_bar(mapping = aes(x = month)) +
  coord_cartesian(ylim = c(2*10^4, 3*10^4))
```



Additional resources:

- r4ds
- dplyr
- dplyr cheat sheet

## Bayesian estimation

We want to minimize with respect to a given loss function -

$$\int L(\hat{\theta} - \theta) * p(\theta|x) d\theta$$

In the lecture, we have seen that when  $L(\hat{\theta}, \theta) = (\hat{\theta} - \theta)^2$  than  $\hat{\theta} = E(\text{posterior})$ , other types of loss functions will derive different estimators (like we have seen above). The logic of this method is as follows - we have somekind of a distribution over  $\theta$ , but we need to choose only one of those. So we formulate an objective function and minimize it with respect to the parameter that we want to find. The most used ones are the mean, median, and common of that distribution, and as we said, they are the bayesian estimators for different loss functions.